


eppstein k-best

knight

Dijkstra



Push A

Pop A & go through A's arcs

Push B.1 (from A)

Push E.50 (from A)

Pop B & go through B's arcs


Push C.3 (from B)

Push D.18 (from B)

Pop C & go through C's arcs

Push Q.7 (from C)

Adjust D.6 (from C)



Analysis:

total pops = n

total pushes = m

maxheapsizes = n


cost of pop = $\log n$

cost of push = 1

$O(m + n \log n)$

K-best

- Eppstein complexity is $O(m + n \log n + k \log k)$
- Great to keep n and k separate – both might be 100,000!
- But consider:



length of 1-best path: n

length of 2-best path: $2n$

...


length of k -best path: kn

total of all these lengths = $n(1 + \dots + k) = O(n k^2)$

- How can Eppstein do better? Well, printing paths is your problem.
- Eppstein computes “implicit representation of k best paths in $O(m + n \log n + k \log k)$ time ... from which we can read off the k th best path in time proportional to the number of edges in that path... and we can read off the length of that path in constant time.”
- *k-best path lengths* problem solvable in $O(m + n \log n + k \log k)$

K-best

- Normal case for speech, MT, and NLG: k-best entries have roughly the same lengths
- Or even:



length of 1-best path: 0

length of 2-best path: 1

length of 3-best path: 1

length of 4-best path: 2

length of 5-best path: 2

length of 6-best path: 2


length of 7-best path: 2

length of 8-best path: 3


length of 9-best path: 3

- Summed lengths of k-best paths = $O(k \log k)$, not k^2 [homework]
- As k goes from 1m to 1b, lengths only increase from 20 to 30

Eppstein






Eppstein




Best path = 55

First Try: Breadth-First Search



Maintain a heap of partial paths.
 Keep popping and pushing.
 When a complete path is popped, print it.
 Vast amount of work even before the 1-best.




First Try: Breadth-First Search



Maintain a heap of partial paths.
Keep popping and pushing.
When a complete path is popped, print it.
Vast amount of work even before the 1-best.



First Try: Breadth-First Search



Maintain a heap of partial paths.

Keep popping and pushing.

When a complete path is popped, print it.


Vast amount of work even before the 1-best.

Core Eppstein idea:

Make a heap where all elements are complete paths.

Every pop gives you the next longest complete path.


Eppstein



Best path = 55

Every node has its own best edge

Eppstein

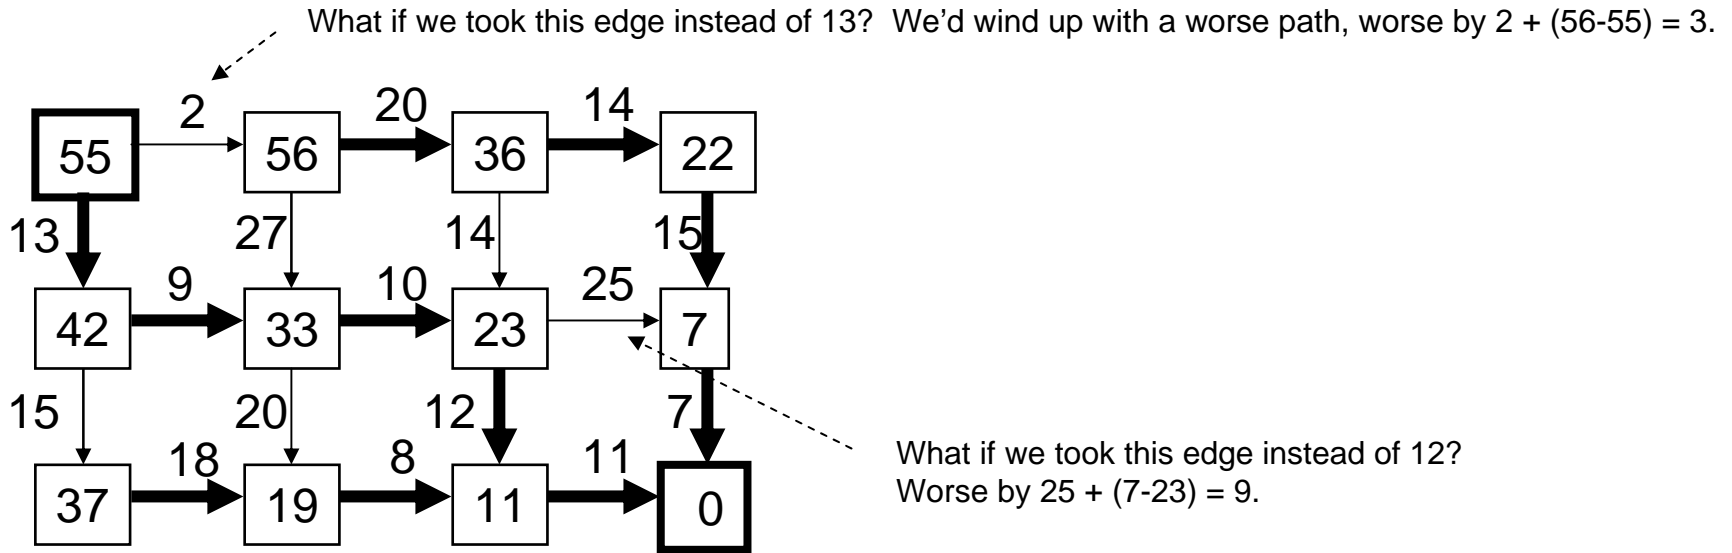


Best path = 55

Every node has its own best edge

Every node has its best cost to goal

Eppstein

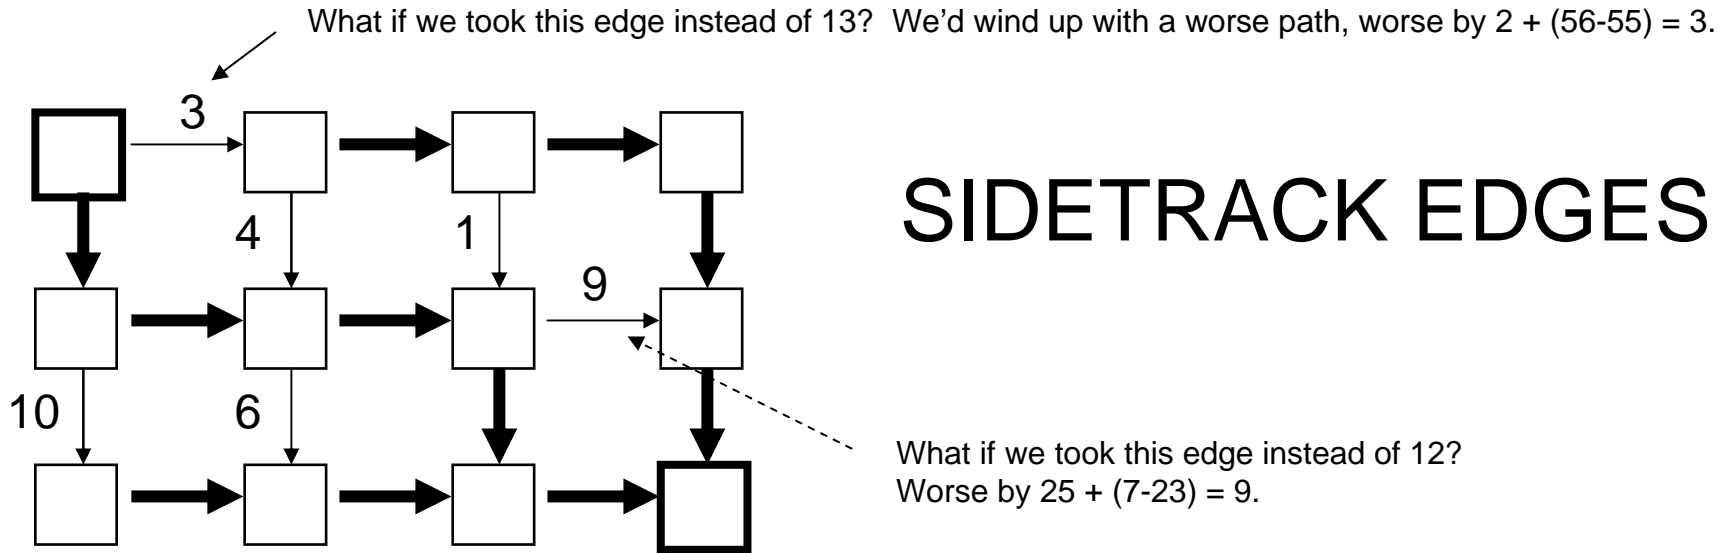


Best path = 55

Every node has its own best edge

Every node has its best cost to goal

Eppstein



Best path = 55


Every node has its own best edge

Every node has its best cost to goal


Every path can be represented as a sequence of sidetrack edges.

So $\{ \}$, $\{3\}$, $\{9\}$, $\{3,1\}$, $\{3, 4, 9\}$, etc. all correspond to paths in original graph.

Eppstein



Best path = 55
 Every node has its own best edge
 Every node has its best cost to goal




Tree of sidetrack sequences
 (Yes, these are all 10 paths!)
 Costs $O(m + n \log n)$ to set this up.


Every path can be represented as a sequence of sidetrack edges.

So $\{\}$, $\{3\}$, $\{9\}$, $\{3,1\}$, $\{3, 4, 9\}$, etc. all correspond to paths in original graph.

Eppstein



Best path = 55





Taking yet each successive sidetrack incurs some incremental cost.

Cost of sidetrack sequence is sum of those incremental costs.

Now, enumerate sidetrack sequences in order of cost.
(& add 55 to each sequence's cost!)

Eppstein



Algorithm (Breadth-First Search):

push root.0 onto heap H

for i = 1 to k



pop node r.c1 from top of H

print cost c1 + 55

for j = each edge from r \rightarrow s with cost c2

push node s.(c1+c2) onto H

Eppstein



Algorithm (Breadth-First Search):

push root.0 onto heap H

for i = 1 to k

pop node r.c1 from top of H



print cost c1 + 55

for j = each edge from r \rightarrow s with cost c2

push node s.(c1+c2) onto H

0.a

Eppstein



55

Algorithm (Breadth-First Search):

push root.0 onto heap H


for i = 1 to k

pop node r.c1 from top of H



print cost c1 + 55

for j = each edge from r → s with cost c2


push node s.(c1+c2) onto H



Eppstein





Algorithm (Breadth-First Search):
 push root.0 onto heap H
 for i = 1 to k
 pop node r.c1 from top of H
 print cost c1 + 55
 for j = each edge from r \rightarrow s with cost c2
 push node s.(c1+c2) onto H



55
58

Eppstein



Algorithm (Breadth-First Search):

push root.0 onto heap H


for i = 1 to k

pop node r.c1 from top of H

print cost c1 + 55



for j = each edge from r → s with cost c2

push node s.(c1+c2) onto H



55
58
59

Eppstein



Algorithm (Breadth-First Search):

push root.0 onto heap H


for $i = 1$ to k

pop node $r.c1$ from top of H

print cost $c1 + 55$



for $j =$ each edge from $r \rightarrow s$ with cost $c2$

push node $s.(c1+c2)$ onto H



55
58
59
61

Eppstein



Algorithm (Breadth-First Search):

push root.0 onto heap H


for i = 1 to k

pop node r.c1 from top of H

print cost c1 + 55



for j = each edge from r → s with cost c2

push node s.(c1+c2) onto H



55
58
59
61
62

Eppstein



Algorithm (Breadth-First Search):

push root.0 onto heap H


for i = 1 to k

pop node r.c1 from top of H

print cost c1 + 55



for j = each edge from r → s with cost c2

push node s.(c1+c2) onto H



55
58
59
61
62
64

Eppstein



Algorithm (Breadth-First Search):

push root.0 onto heap H


for i = 1 to k

pop node r.c1 from top of H

print cost c1 + 55



for j = each edge from r → s with cost c2

push node s.(c1+c2) onto H



55
58
59
61
62
64
65

Eppstein



Algorithm (Breadth-First Search):

push root.0 onto heap H


for i = 1 to k

pop node r.c1 from top of H

print cost c1 + 55



for j = each edge from r \rightarrow s with cost c2

push node s.(c1+c2) onto H



55
58
59
61
62
64
65
68

Eppstein



Algorithm (Breadth-First Search):

push root.0 onto heap H

for i = 1 to k

pop node r.c1 from top of H

print cost c1 + 55



for j = each edge from r → s with cost c2

push node s.(c1+c2) onto H

16.i

55
58
59
61
62
64
65
68
68

Eppstein





Algorithm (Breadth-First Search):
 push root.0 onto heap H
 for i = 1 to k
 pop node r.c1 from top of H
 print cost c1 + 55
 for j = each edge from r → s with cost c2
 push node s.(c1+c2) onto H

<empty heap>

55
 58
 59
 61
 62
 64
 65
 68
 68
 71

Eppstein



Analysis:

total pops = k

total pushes = km (if fully connected)

maxheapsizes = $O(km)$

cost of pop = $O(\log km) = O(\log k + \log m)$

cost of push = 1

→ $O(km + k \log k + k \log m)$

→ $O(km + k \log k)$

→ $km = kn^2 = \text{bad term to have!!}$

Algorithm (Breadth-First Search):

push root.0 onto heap H

for $i = 1$ to k



pop node $r.c1$ from top of H

print cost $c1 + 55$

for $j = \text{each edge from } r \rightarrow s \text{ with cost } c2$

push node $s.(c1+c2)$ onto H

Eppstein



Algorithm (Breadth-First Search):
 push root.0 onto heap H
 for i = 1 to k
 pop node r.c1 from top of H
 print cost c1 + 55
 for j = each edge from r \rightarrow s with cost c2
 push node s.(c1+c2) onto H

Analysis for bounded outdegree tree:

total pops = k

total pushes = k * constant

maxheapsizes = $O(k)$



cost of pop = $O(\log k)$

cost of push = 1

$\rightarrow O(k + k \log k)$



$\rightarrow O(k \log k)$, no n factor

Eppstein





Algorithm (Breadth-First Search):
 push root.0 onto heap H
 for i = 1 to k
 pop node r.c1 from top of H
 print cost c1 + 55
 for j = each edge from r → s with cost c2
 push node s.(c1+c2) onto H

Eppstein





Algorithm (Breadth-First Search):
 push root.0 onto heap H
 for i = 1 to k
 pop node r.c1 from top of H
 print cost c1 + 55
 for j = each edge from r \rightarrow s with cost c2
 push node s.(c1+c2) onto H

Eppstein





Algorithm (Breadth-First Search):
 push root.0 onto heap H
 for i = 1 to k
 pop node r.c1 from top of H
 print cost c1 + 55
 for j = each edge from r \rightarrow s with cost c2
 push node s.(c1+c2) onto H

Eppstein



Algorithm (Breadth-First Search):
 push root.0 onto heap H
 for i = 1 to k
 pop node r.c1 from top of H
 print cost c1 + 55
 for j = each edge from r \rightarrow s with cost c2
 push node s.(c1+c2) onto H

Eppstein




Algorithm (Breadth-First Search):
 push root.0 onto heap H
 for i = 1 to k
 pop node r.c1 from top of H
 print cost c1 + 55
 for j = each edge from r \rightarrow s with cost c2
 push node s.(c1+c2) onto H

So we might be able to
 reduce the outdegree...





Goals:


- bounded outdegree sidetrack tree
- maintain 1-to-1 correspondence of
 - sidetrack tree paths
 - original graph start \rightarrow final paths



1. Pick best sidetrack of each node from start to final. Arrange all in binary heap.



2. Repeat for other nodes besides start.








1. Pick best sidetrack of each node from start to final. Arrange all in binary heap.

2. Repeat for other nodes besides start.





3. Record incremental costs.





1. Pick best sidetrack of each node from start to final. Arrange all in binary heap.



2. Repeat for other nodes besides start.



CROSS
EDGES


3. Record incremental costs.

4. Connect each sidetrack node to the heap of its tail. Outdegree = 3.




1. Pick best sidetrack of each node from start to final. Arrange all in binary heap.

2. Repeat for other nodes besides start.






3. Record incremental costs.






4. Connect each sidetrack node to the heap of its tail. Outdegree = 3.

CROSS EDGES

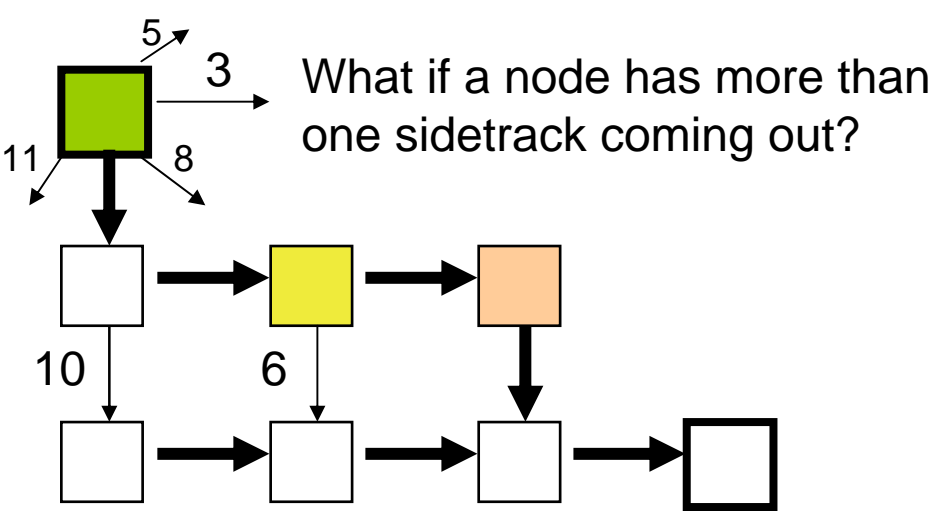
5. Erase node labels & add special start "+0" node.






This new tree has the same 10 paths!
 But it has max outdegree = 3.
 Now same algorithm (BFS) works better!

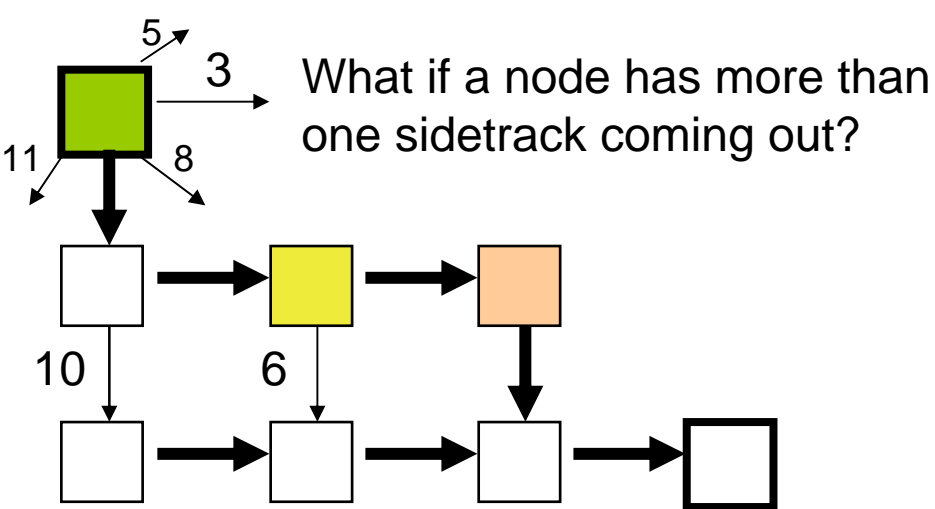
DONE!



Detail #1


multiple sidetracks





Detail #1

multiple sidetracks




Detail #2

structure sharing

Notice that 6 and 9 appear many times.

Heaps need to be built so that they share a bunch of sub-structure.


Done carefully, the whole structure can be built in $O(m + n \log n)$ time, leaving $O(k \log k)$ for extracting the items using BFS.




Detail #3

loops are no problem

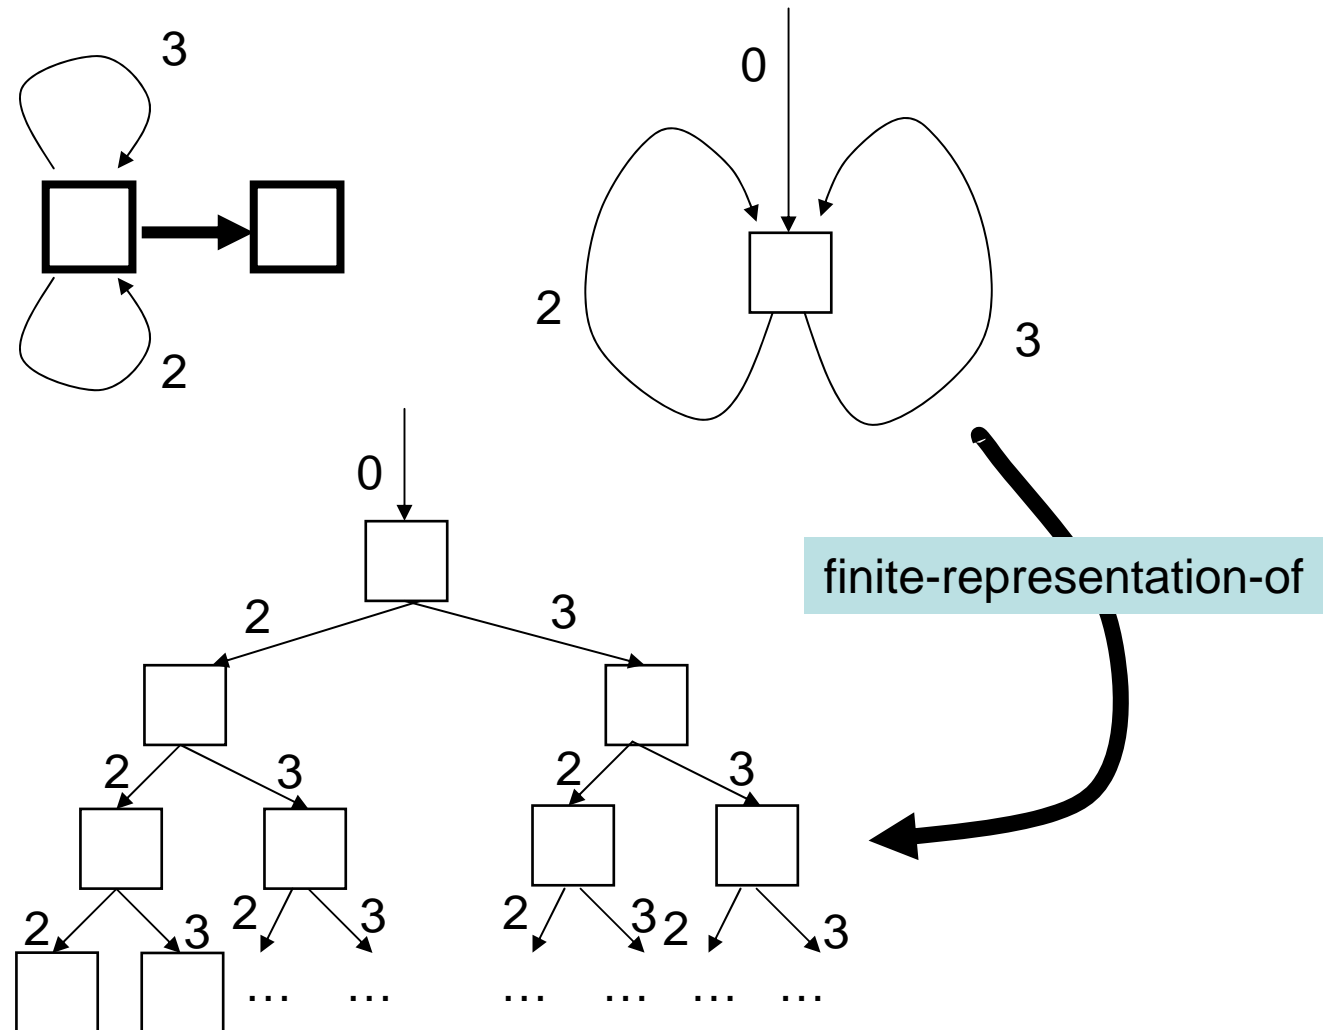
Original graph




Sidetrack edges



Tree of sidetrack sequences



K-Best Dijkstra v. Eppstein



K-Best Dijkstra

Push (S,0)

Pop (S,0)

Do m+1 pushes

Pop (S,1)

Do m+1 pushes


Pop (S,1)

Do m+1 pushes

...

At least $O(km)$ work
before any complete
path is generated...

Eppstein



Each BFS-pop corresponds to
a complete path:

1-1000, 2-1000, 1-1-1000, 3-1000, 1-2-
1000, 2-1-1000, 1-1-1-1000,
4-1000, 3-1-1000, 1-3-1000, 2-2-1000, 1-
1-2-1000, ...

the end